

Devon Perez

New York, NY • (468) 502-3216 • devonperez647@proton.me

LinkedIn: [linkedin.com/in/devonperez647](https://www.linkedin.com/in/devonperez647) • GitHub: github.com/devonperez647

SUMMARY

Junior ml platform engineer with 3+ years experience across CV, MLOps. Focus on measurable business impact, production reliability, and fast iteration.

SKILLS

A/B testing, Bandits, CNNs, CUDA, Docker, Hugging Face, JAX, Kafka, Kubernetes, LangChain, MLflow, Redis, XGBoost, YOLO

EXPERIENCE

ML Platform Engineer, VectorWorks – 2021–2024

- Built and shipped ML systems end-to-end (A/B testing, Redis, Bandits, Hugging Face, Kafka, CNNs).
- Collaborated with product/infra; mentored 4 engineers.

PROJECTS

- Vector Search (RecSys): Built Faiss pipeline for candidate gen, 50ms P95 for 50M vectors.
- Personalized Recommendations (RecSys): Hybrid deep+wide model; +11% CTR and +7% revenue per session in A/B.
- Defect Detection (CV): Deployed real-time vision model on edge GPUs; reduced false negatives by 35% at 60 FPS.
- OCR Pipeline (CV): Trained transformer-based OCR; boosted accuracy on noisy scans from 82%→94%.

EDUCATION

B.S. Computer Science, University of Illinois Urbana-Champaign (2016)

CERTIFICATIONS

TensorFlow Developer Certificate