

Hayden Hernandez
Los Angeles, CA • (606) 477-6268 • haydenhernandez972@gmail.com
LinkedIn: [linkedin.com/in/haydenhernandez972](https://www.linkedin.com/in/haydenhernandez972) • GitHub:
github.com/haydenhernandez972

SUMMARY

Mid-level mlops engineer with 6+ years experience across RecSys, TimeSeries. Focus on measurable business impact, production reliability, and fast iteration.

SKILLS

A/B testing, CUDA, Docker, Faiss, Kafka, Matrix Factorization, ONNX, PostgreSQL, ROCm, Segmentation, Triton kernels, Vector DBs, Weights & Biases, YOLO

EXPERIENCE

MLOps Engineer, Lattice Solutions – 2018–2022

- Built and shipped ML systems end-to-end (PostgreSQL, Triton kernels, Kafka, Weights & Biases, ONNX, YOLO).
- Collaborated with product/infra; mentored 2 engineers.

MLOps Engineer, VectorWorks – 2014–2018

- Built and shipped ML systems end-to-end (Segmentation, Docker, Triton kernels, Matrix Factorization, ONNX, Weights & Biases).
- Collaborated with product/infra; mentored 1 engineers.

PROJECTS

- Vector Search (RecSys): Built Faiss pipeline for candidate gen, 50ms P95 for 50M vectors.
- Personalized Recommendations (RecSys): Hybrid deep+wide model; +11% CTR and +7% revenue per session in A/B.
- Defect Detection (CV): Deployed real-time vision model on edge GPUs; reduced false negatives by 35% at 60 FPS.
- OCR Pipeline (CV): Trained transformer-based OCR; boosted accuracy on noisy scans from 82%→94%.

EDUCATION

M.S. Computer Science, Cornell University (2018)

CERTIFICATIONS

TensorFlow Developer Certificate, Google Professional ML Engineer