

Kai Martinez

New York, NY • (250) 977-5584 • kaimartinez336@proton.me

LinkedIn: [linkedin.com/in/kaimartinez336](https://www.linkedin.com/in/kaimartinez336) • GitHub: github.com/kaimartinez336

SUMMARY

Mid-level mlops engineer with 3+ years experience across MLOps, RecSys. Focus on measurable business impact, production reliability, and fast iteration.

SKILLS

Airflow, BERT, Batch & Streaming, CI/CD, Faiss, JAX, LangChain, MLflow, PostgreSQL, Python, ROCm, Spacy, TensorFlow, Triton kernels

EXPERIENCE

MLOps Engineer, AcmeAI – 2018–2021

- Built and shipped ML systems end-to-end (LangChain, ROCm, Airflow, Triton kernels, PostgreSQL, Batch & Streaming).
- Collaborated with product/infra; mentored 3 engineers.

PROJECTS

- Model Observability (MLOps): Implemented drift monitoring; MTTR down from days to hours.
- Feature Store (MLOps): Designed Feast-based store; cut feature duplication by 70%.
- Domain-specific RAG Assistant (NLP): Built retrieval-augmented LLM for internal knowledge base; cut search time by 62% and enabled 5k daily queries.
- LLM Fine-tuning (NLP): Fine-tuned 7B model with LoRA+QLoRA; improved factual QA F1 from 63→78 on domain dataset of 80k pairs.

EDUCATION

Ph.D. Computer Science, University of Washington (2022)

CERTIFICATIONS

AWS Certified ML – Specialty, Google Professional ML Engineer