

Devon Martinez
Atlanta, GA • (475) 687-5463 • devonmartinez694@gmail.com
LinkedIn: linkedin.com/in/devonmartinez694 • GitHub:
github.com/devonmartinez694

SUMMARY

Mid-level ml platform engineer with 3+ years experience across RecSys, TimeSeries. Focus on measurable business impact, production reliability, and fast iteration.

SKILLS

ARIMA, Beam search, DVC, Docker, Faiss, Kafka, LLMs, LSTMs, LangChain, LightGBM, MLflow, ROCm, Redis, SQL

EXPERIENCE

ML Platform Engineer, Pixeltree – 2017–2020

- Built and shipped ML systems end-to-end (LLMs, LightGBM, Faiss, Beam search, Kafka, Docker).
- Collaborated with product/infra; mentored 5 engineers.

PROJECTS

- Anomaly Detection (TimeSeries): Streaming detector with isolation forests; reduced alert fatigue 40%.
- Forecasting Platform (TimeSeries): Auto-ML for SKU forecasting; MAPE down 18% across 1.2k series.
- LLM Fine-tuning (NLP): Fine-tuned 7B model with LoRA+QLoRA; improved factual QA F1 from 63→78 on domain dataset of 80k pairs.
- Domain-specific RAG Assistant (NLP): Built retrieval-augmented LLM for internal knowledge base; cut search time by 62% and enabled 5k daily queries.

EDUCATION

B.S. Electrical Engineering, University of Washington (2019)

CERTIFICATIONS

—