

Jules Garcia
San Francisco, CA • (485) 503-4009 • julesgarcia182@gmail.com
LinkedIn: [linkedin.com/in/julesgarcia182](https://www.linkedin.com/in/julesgarcia182) • GitHub: github.com/julesgarcia182

SUMMARY

Mid-level senior machine learning engineer with 5+ years experience across NLP, Tabular. Focus on measurable business impact, production reliability, and fast iteration.

SKILLS

Airflow, Bandits, Hugging Face, Implicit feedback, Kafka, LLMs, LangChain, LightGBM, MLflow, PyTorch, Python, Redis, Spacy, Weights & Biases

EXPERIENCE

Senior Machine Learning Engineer, Nimbus Labs – 2014–2018

- Built and shipped ML systems end-to-end (Spacy, MLflow, Airflow, LangChain, Hugging Face, Weights & Biases).
- Collaborated with product/infra; mentored 2 engineers.

PROJECTS

- Vector Search (RecSys): Built Faiss pipeline for candidate gen, 50ms P95 for 50M vectors.
- Personalized Recommendations (RecSys): Hybrid deep+wide model; +11% CTR and +7% revenue per session in A/B.
- Domain-specific RAG Assistant (NLP): Built retrieval-augmented LLM for internal knowledge base; cut search time by 62% and enabled 5k daily queries.
- LLM Fine-tuning (NLP): Fine-tuned 7B model with LoRA+QLoRA; improved factual QA F1 from 63→78 on domain dataset of 80k pairs.

EDUCATION

B.S. Computer Science, Stanford University (2012)

CERTIFICATIONS

—