

Quinn Taylor

Los Angeles, CA • (626) 848-5130 • quinntaylor733@gmail.com

LinkedIn: [linkedin.com/in/quinntaylor733](https://www.linkedin.com/in/quinntaylor733) • GitHub: github.com/quinntaylor733

SUMMARY

Senior ml platform engineer with 8+ years experience across RecSys, CV. Focus on measurable business impact, production reliability, and fast iteration.

SKILLS

ANN search, Docker, Faiss, Kubernetes, LightGBM, Matrix Factorization, Model Registry, ONNX, Observability, PyTorch, Ray, TensorFlow, Vector DBs, scikit-learn

EXPERIENCE

ML Platform Engineer, Nimbus Labs – 2020–2024

- Built and shipped ML systems end-to-end (Matrix Factorization, TensorFlow, Kubernetes, ONNX, Model Registry, Ray).
- Collaborated with product/infra; mentored 3 engineers.

ML Platform Engineer, NeuroNova – 2012–2016

- Built and shipped ML systems end-to-end (Kubernetes, LightGBM, Matrix Factorization, PyTorch, Faiss, ONNX).
- Collaborated with product/infra; mentored 3 engineers.

PROJECTS

- Model Observability (MLOps): Implemented drift monitoring; MTTR down from days to hours.
- Feature Store (MLOps): Designed Feast-based store; cut feature duplication by 70%.
- Vector Search (RecSys): Built Faiss pipeline for candidate gen, 50ms P95 for 50M vectors.
- Personalized Recommendations (RecSys): Hybrid deep+wide model; +11% CTR and +7% revenue per session in A/B.

EDUCATION

Ph.D. Computer Science, Cornell University (2018)

CERTIFICATIONS

Databricks ML Practitioner