

Quinn Chen
Seattle, WA • (506) 697-2905 • quinnchen148@proton.me
LinkedIn: [linkedin.com/in/quinnchen148](https://www.linkedin.com/in/quinnchen148) • GitHub: github.com/quinnchen148

SUMMARY

Mid-level ml platform engineer with 3+ years experience across Tabular, MLOps. Focus on measurable business impact, production reliability, and fast iteration.

SKILLS

Airflow, Bandits, Faiss, Kalman Filters, LSTMs, LangChain, LightGBM, Matrix Factorization, PostgreSQL, SQL, Triton kernels, Vector DBs, Weights & Biases, scikit-learn

EXPERIENCE

ML Platform Engineer, BlueWave Analytics – 2015–2018

- Built and shipped ML systems end-to-end (SQL, Airflow, Matrix Factorization, LangChain, Vector DBs, Weights & Biases).
- Collaborated with product/infra; mentored 5 engineers.

PROJECTS

- Vector Search (RecSys): Built Faiss pipeline for candidate gen, 50ms P95 for 50M vectors.
- Personalized Recommendations (RecSys): Hybrid deep+wide model; +11% CTR and +7% revenue per session in A/B.
- Anomaly Detection (TimeSeries): Streaming detector with isolation forests; reduced alert fatigue 40%.
- Forecasting Platform (TimeSeries): Auto-ML for SKU forecasting; MAPE down 18% across 1.2k series.

EDUCATION

Ph.D. Computer Science, Harvard University (2014)

CERTIFICATIONS

Azure AI Engineer Associate, NVIDIA Deep Learning Institute